

P8130 HW3

10/16/2020

Problem 1

```
# data import
exercise = read_csv(file = "./data/Exercise.csv") %>%
  janitor::clean_names() %>%
  select(group, systolic_pre, systolic_post) %>%
  mutate(systolic_diff = systolic_post - systolic_pre) # difference
```

```
## Parsed with column specification:
## cols(
##   Group = col_double(),
##   Age = col_double(),
##   Gender = col_double(),
##   Race = col_double(),
##   HTN = col_double(),
##   T2DM = col_double(),
##   Depression = col_double(),
##   Smokes = col_double(),
##   Systolic_PRE = col_double(),
##   Systolic_POST = col_double()
## )
```

(a) Since we are comparing the changes of Systolic BP for same patients at two different timepoints, We will employ ‘Two-Sided Paired t-test’ to assess whether the Systolic BP at 6 months is significantly from the baseline values for each of the groups: intervention, control. (assumptions are checked in part (c)) *intervention group*

$$H_0 : \mu_{post} - \mu_{pre} = 0$$

$$H_1 : \mu_{post} - \mu_{pre} \neq 0$$

With a pre-specified significance level $\alpha = 0.05$, the test statistics

$$t_{stat} = \frac{\bar{d} - 0}{s_d / \sqrt{n}}$$

Reject H_0 : if $|t_{stat}| > t_{n-1, 1-\alpha/2}$

Fail to reject H_0 : if $|t_{stat}| \leq t_{n-1, 1-\alpha/2}$

We find the following:

$$\bar{d} = -8.583333$$

$$s_d = 17.1687$$

$$n = 36$$

$$t_{stat} = \frac{\bar{d}-0}{s_d/\sqrt{n}} = \frac{-8.583333-0}{17.1687/\sqrt{36}} \cong -2.999645$$

$$t_{crit} = qt(0.975, 35) = 2.030108$$

- $|t_{stat}| > t_{35, 0.975}$
- Conclusion: Reject H_0 , and conclude that the Systolic BP at 6 month is significant different from the baseline values for the intervention group @ $\alpha = 0.05$.

control group

$$H_0 : \mu_{post} - \mu_{pre} = 0$$

$$H_1 : \mu_{post} - \mu_{pre} \neq 0$$

With a pre-specified significance level $\alpha = 0.05$, the test statistics

$$t_{stat} = \frac{\bar{d}-0}{s_d/\sqrt{n}}$$

Reject H_0 : if $|t_{stat}| > t_{n-1, 1-\alpha/2}$

Fail to reject H_0 : if $|t_{stat}| \leq t_{n-1, 1-\alpha/2}$

We find the following:

$$\bar{d} = -3.333333$$

$$s_d = 14.81312$$

$$n = 36$$

$$t_{stat} = \frac{\bar{d}-0}{s_d/\sqrt{n}} = \frac{-3.333333-0}{14.81312/\sqrt{36}} \cong -1.350154$$

$$t_{crit} = qt(0.975, 35) = 2.030108$$

- $|t_{stat}| < t_{35, 0.975}$
- Conclusion: Fail to reject H_0 , and conclude that the Systolic BP at 6 month is not significant different from the baseline values for the control group @ $\alpha = 0.05$.

```
## intervention group:
intervention = exercise %>% filter(group == 1)
mean_diff_int = mean(pull(intervention, systolic_diff)) # -8.583333
sd_diff_int = sd(pull(intervention, systolic_diff)) #17.1687
n_int = nrow(intervention) #36
t_stat_int = (mean_diff_int - 0)/(sd_diff_int/sqrt(n_int)) #-2.999645
t_crit = qt(0.975, 35) #2.030108
## summary:
t.test(intervention$systolic_post, intervention$systolic_pre, paired = T)

##
## Paired t-test
##
## data: intervention$systolic_post and intervention$systolic_pre
## t = -2.9996, df = 35, p-value = 0.004953
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -14.392386 -2.774281
## sample estimates:
## mean of the differences
## -8.583333
```

```
## control group:
control = exercise %>% filter(group == 0)
mean_diff_con = mean(control$systolic_diff) # -3.333333
sd_diff_con = sd(control$systolic_diff) #14.81312
n_con = nrow(control) #36
t_stat_con = (mean_diff_con - 0)/(sd_diff_con/sqrt(n_con)) #-1.350154
t_crit = qt(0.975, 35) #2.030108
## summary:
t.test(control$systolic_post, control$systolic_pre, paired = T)

##
## Paired t-test
##
## data: control$systolic_post and control$systolic_pre
## t = -1.3502, df = 35, p-value = 0.1856
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -8.345373 1.678706
## sample estimates:
## mean of the differences
## -3.333333
```

(b) Since patients in the two groups are totally different, we will assess the Systolic BP absolute changes between the two groups by using an ‘Two-sample Test for Independent samples’ Then, the first step is to check whether the underlying variances of the two groups are equal or not. In this case, yes. Test for Equality of Variance

$$H_0 : \sigma_{intervention}^2 = \sigma_{control}^2$$

$$H_1 : \sigma_{intervention}^2 \neq \sigma_{control}^2$$

Compute the F test statistics and critical value: $F_{stat} = \frac{s_{intervention}^2}{s_{control}^2} = \frac{17.1687}{14.81312} \cong 1.3433268$

$$F_{crit} = \text{qf}(0.975, 35, 35) = 1.9610894$$

- Decision: At 0.05 significance level, because $F_{stat} < F_{crit}$, we fail to reject H_0 and conclude that it is safe to assume variances are equal.

The pooled estimate of the variance from the two independent groups is given by:

$$s^2 = \frac{(n_{int}-1)s_{int}^2 + (n_{con}-1)s_{con}^2}{n_{int}+n_{con}-2} = \frac{(36-1)17.1687 + (36-1)14.81312}{36+36-2} \cong 257.0964 \text{ Where } s = \sqrt{s^2} = 16.03423$$

Therefore, we have the following hypothesis test:

$$H_0 : \mu_{intervention} = \mu_{control}$$

$$H_1 : \mu_{intervention} \neq \mu_{control}$$

$$t_{stat} = \frac{\bar{\mu}_{intervention} - \bar{\mu}_{intervention}}{s/\sqrt{1/n_1 + 1/n_2}} \sim t_{n_1+n_2-2} \text{ under the } H_0$$

$$t_{stat} = \frac{(-8.583333 - (-3.333333))}{16.03423/\sqrt{1/36 + 1/36}} \cong -1.389145$$

$$t_{crit} = \text{qt}(0.975, 70) = 1.994437 > |t_{stat}| = 1.389145$$

- Conclusion: Fail to reject H_0 , and conclude the difference in mean systolic blood pressure change (post-pre) is not significantly different between the intervention group and control group @ $\alpha = 0.05$.

95% Confidence Interval:

$$\bar{X}_{intervention} - \bar{X}_{control} \pm t_{n_{intervention}+n_{control}-2, 1-\alpha/2} \times \frac{s}{\sqrt{1/n_1+1/n_2}}$$
$$(-8.583333 - (-3.333333)) \pm 1.995469 \times \sqrt{\frac{257.0964}{2/72}}$$
$$(-197.1255, 186.6255)$$

- Conclusion: we are 95% confident that the true Systolic BP absolute changes for the two groups is between -12.78758 and 2.287583.

```
# variance test #
var.test(intervention$systolic_diff, control$systolic_diff, alternative = "two.sided")

##
## F test to compare two variances
##
## data: intervention$systolic_diff and control$systolic_diff
## F = 1.3433, num df = 35, denom df = 35, p-value = 0.3869
## alternative hypothesis: true ratio of variances is not equal to 1
## 95 percent confidence interval:
## 0.6849901 2.6343840
## sample estimates:
## ratio of variances
## 1.343327

## Alternatively:
# F_stat
F_stat = (sd_diff_int)^2/(sd_diff_con)^2#1.343327
# F_crit
F_crit = qf(0.975,35,35)#1.961089
#####

# Independent two samples test #
t.test(intervention$systolic_diff, control$systolic_diff, var.equal = TRUE, paired = FALSE)

##
## Two Sample t-test
##
## data: intervention$systolic_diff and control$systolic_diff
## t = -1.3891, df = 70, p-value = 0.1692
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -12.787583 2.287583
## sample estimates:
## mean of x mean of y
## -8.583333 -3.333333

## Alternatively:
# s^2, s
s_2 = (35*sd_diff_int^2 + 35*sd_diff_con^2)/(36 + 36 - 2)#257.0964
s = sqrt(s_2)#16.03423
# test statistics
(mean_diff_int - mean_diff_con)/(s*sqrt(1/36 + 1/36))# |-1.389145| --> 1.389145
```

```
## [1] -1.389145
```

```
# test critical  
qt(0.975,70)#1.994437 > test statistics --> fail to reject H0
```

```
## [1] 1.994437
```

```
# 95% C.I  
mean_diff_int - mean_diff_con #-5.25
```

```
## [1] -5.25
```

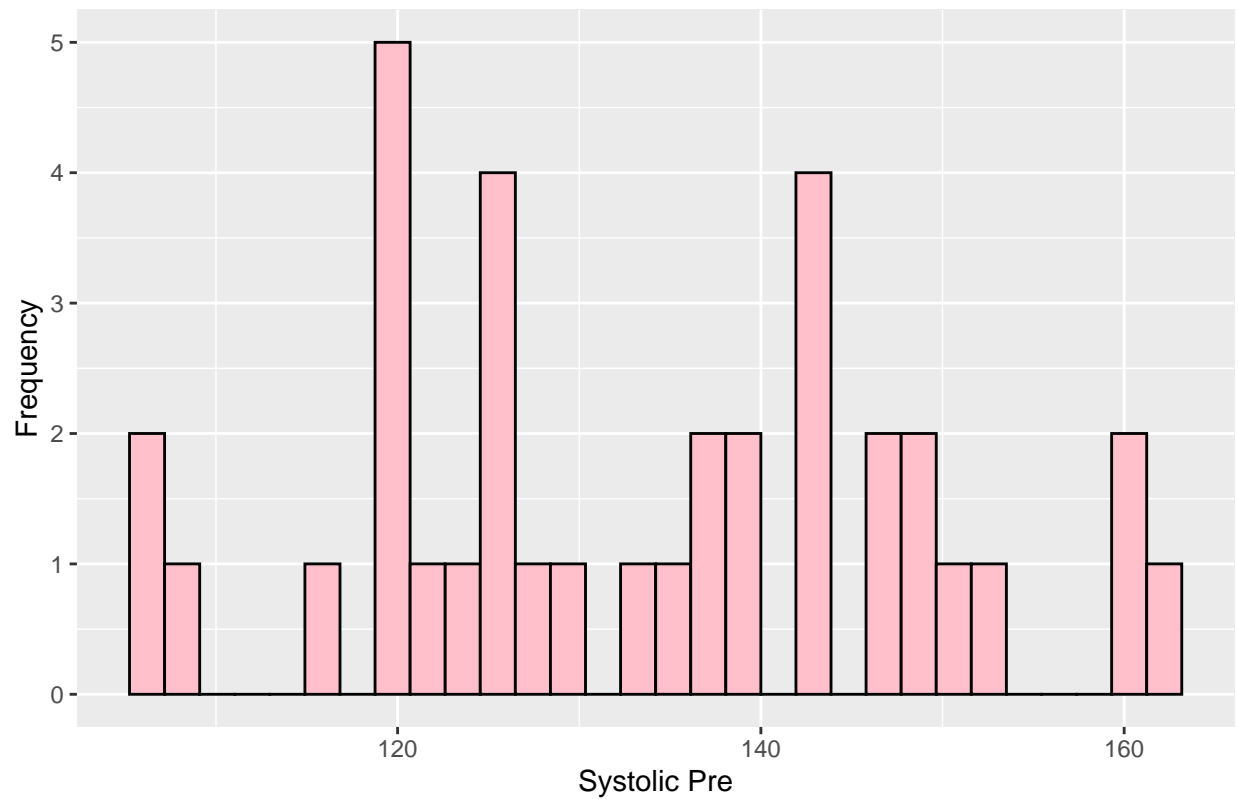
```
lower_bound = (mean_diff_int - mean_diff_con) - qt(0.975,70)*s*sqrt(2/36)#-12.78758  
upper_bound = (mean_diff_int - mean_diff_con) + qt(0.975,70)*s*sqrt(2/36)#2.287583
```

(c) What are the main underlying assumptions for the tests performed in parts a) and b)?*

- By CLT if you are sampled from a underlying normal distribution, the sample mean also follows a normal distribution, and the difference if the sample means are also follows a normal distribution.
- For part a), we use a paired t-test, and we assume the SBP measurements(post and pre) for both intervention and control groups are normally distributed with mean $\mu_{int_{post}}$ and $\mu_{int_{pre}}$, and $\mu_{con_{post}}$ and $\mu_{con_{pre}}$, respectively. It follows that the differences $d_{i_{int}} \sim N(\Delta_{int}, \sigma_{d_{int}}^2)$, $d_{i_{con}} \sim N(\Delta_{con}, \sigma_{d_{con}}^2)$, where $i=1,2,3,\dots,n$. Under the null hypothesis, the test statistics follows a t_{n-1} .
- For part b), we consider the the two groups are independent samples, and that they are normally distributed: $X_{intervention_{\Delta}} \sim N(\mu_{int_{\Delta}}, \sigma_{int_{\Delta}}^2)$ and $X_{control_{\Delta}} \sim N(\mu_{int_{\Delta}}, \sigma_{int_{\Delta}}^2)$. And we found the underlying variances of the two samples are equal.

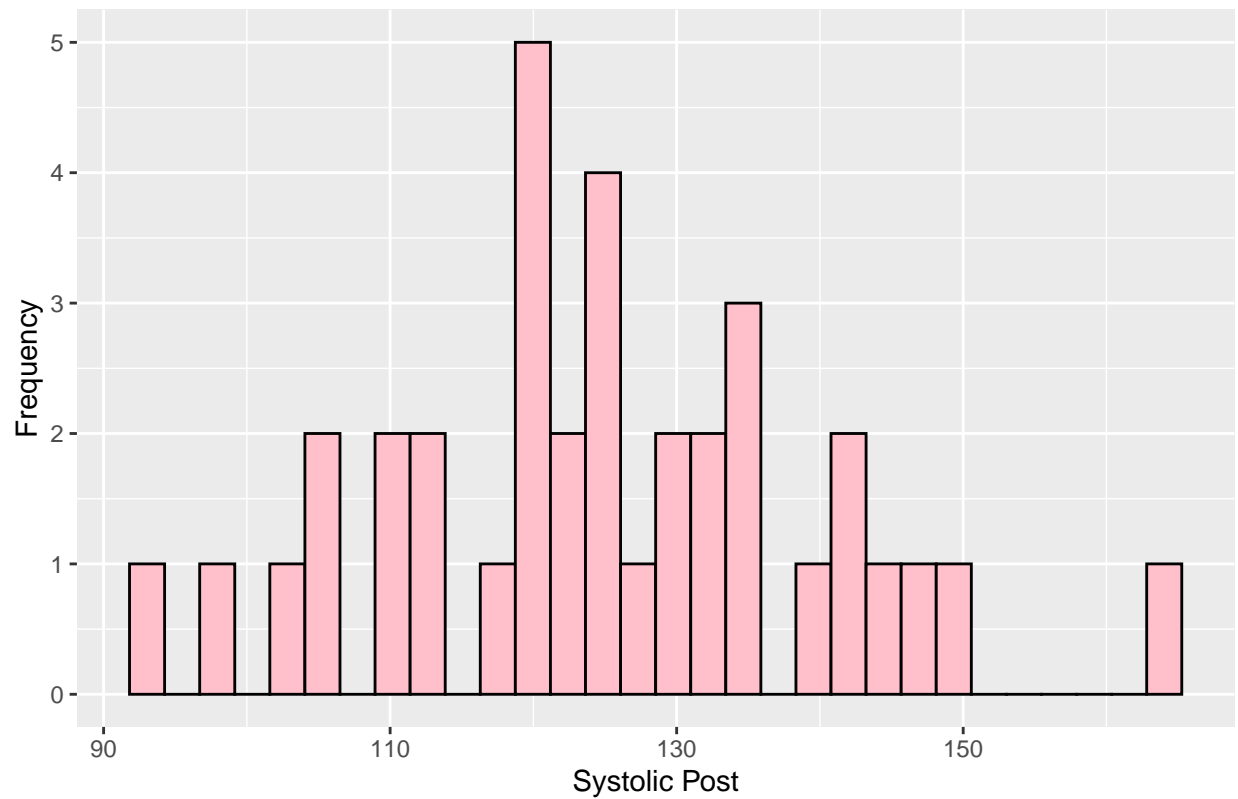
```
#normality_intervention_pre  
par(mfrow = c(2,2))  
ggplot(intervention, aes(x = systolic_pre)) +  
  geom_histogram(color = "black", fill = "pink", bins = 30) +  
  labs(  
    title = "The distribution of Systolic_pre of intervention group",  
    x = "Systolic Pre",  
    y = "Frequency"  
  )
```

The distribution of Systolic_pre of intervention group




```
#normality_intervention_post
ggplot(intervention, aes(x = systolic_post)) +
  geom_histogram(color = "black", fill = "pink", bins = 30) +
  labs(
    title = "The distribution of Systolic_post of intervention group",
    x = "Systolic Post",
    y = "Frequency"
  )
```

The distribution of Systolic_post of intervention group




```
#normality_control_pre
ggplot(control, aes(x = systolic_pre)) +
  geom_histogram(color = "black", fill = "pink", bins = 30) +
  labs(
    title = "The distribution of Systolic_pre of control group",
    x = "Systolic Pre",
    y = "Frequency"
  )
```

The distribution of Systolic_pre of control group



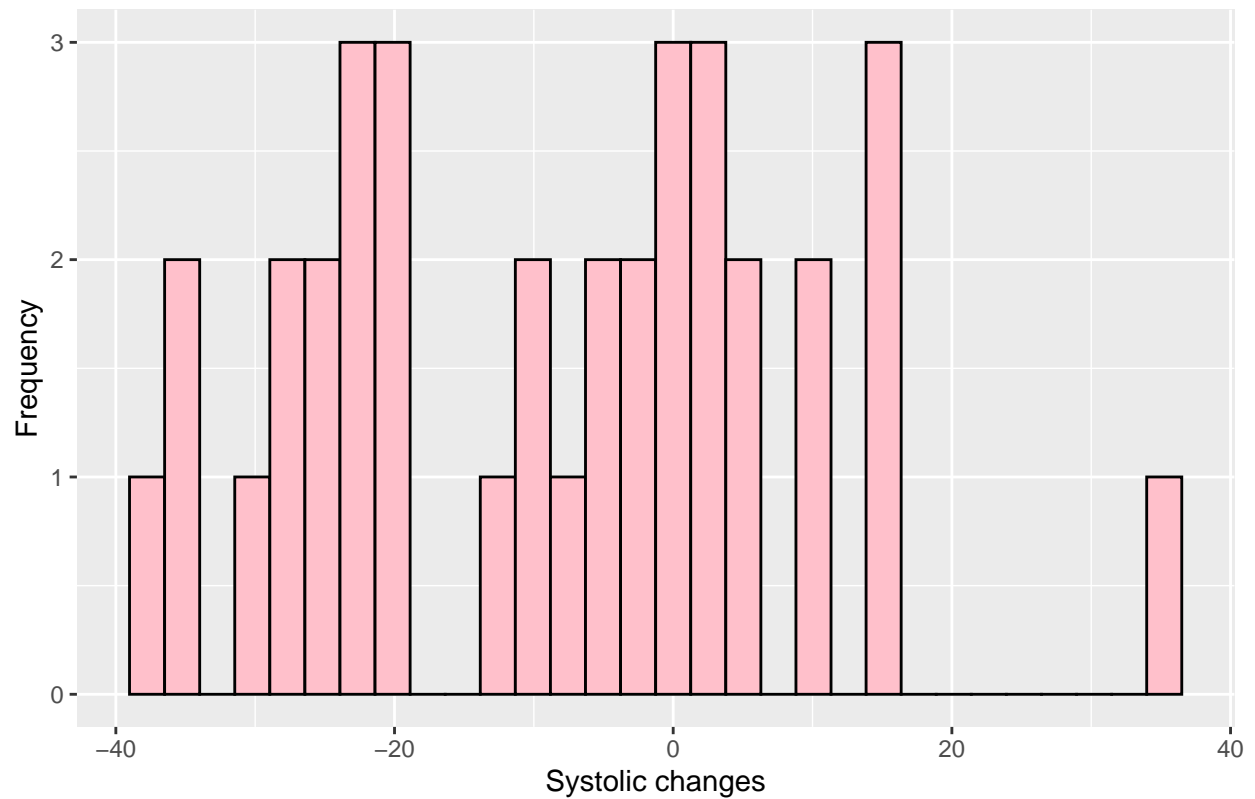
```
#normality_control_pre
ggplot(control, aes(x = systolic_post)) +
  geom_histogram(color = "black", fill = "pink", bins = 30) +
  labs(
    title = "The distribution of Systolic_post of control group",
    x = "Systolic Post",
    y = "Frequency"
  )
```


The distribution of Systolic_post of control group

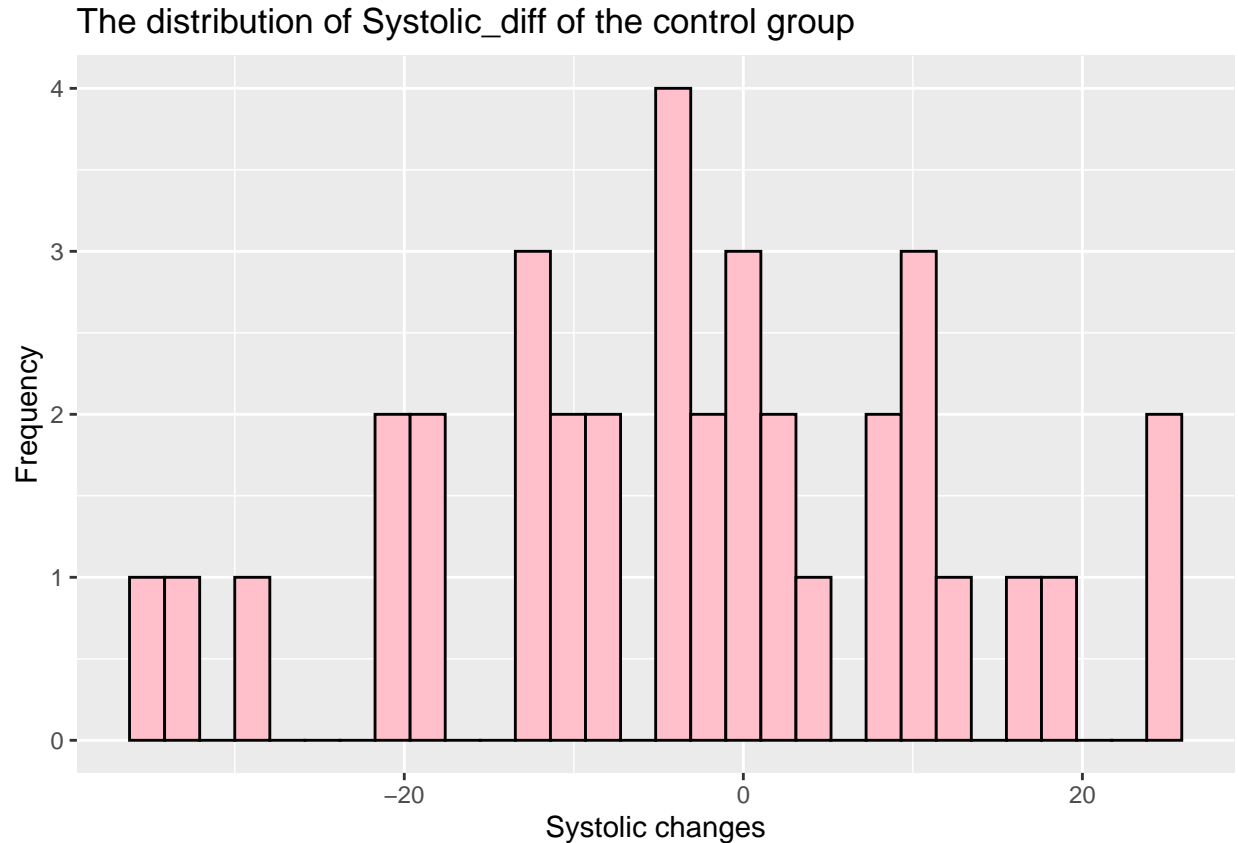


```
#normality_intervention_diff
ggplot(intervention, aes(x = systolic_diff)) +
  geom_histogram(color = "black", fill = "pink", bins = 30) +
  labs(
    title = "The distribution of Systolic_diff of the intervention group",
    x = "Systolic changes",
    y = "Frequency"
  )
```

The distribution of Systolic_diff of the intervention group



```
#normality_control_diff
ggplot(control, aes(x = systolic_diff)) +
  geom_histogram(color = "black", fill = "pink", bins = 30) +
  labs(
    title = "The distribution of Systolic_diff of the control group",
    x = "Systolic changes",
    y = "Frequency"
  )
```



- By visualizing the above histograms, it's hard to state the assumptions of normality are valid since they do not quite look like a symmetric bell-curve. It means the tests we employed above might not be appropriate. Possible remedies to this situation would be increase the total sample sizes for both groups. Therefore, with a large sample size, the normality assumption holds approximately due to CLT.

Problem 2

One_tailed Hypothesis test:

$$H_0 : \mu = \mu_0 = 120 \text{ v.s. } H_1 : \mu < 120$$

Where $\sigma = 15$, $\alpha = 0.05$

$$Z_{stat} = \frac{\bar{X} - \mu_0}{\sigma / \sqrt{n}}$$

$$Z_{crit} = \text{qnorm}(0.05) \cong -1.65$$

Reject H_0 : if $Z_{stat} < Z_{crit}$ \$ Fail to reject H_0 : if $Z_{stat} \geq Z_{crit}$

`rnorm(n = 20, mean = 120, sd = 15)`

(a) **one random sample from the underlying (null) true distribution** We find: $|Z_{stat}| = 0.01776018 > Z_{crit} = -1.644854$, fail to reject H_0 , and conclude that the true mean of average IQ score of IVY League is 270. The conclusion is 1 in this one random sample case.

```
set.seed(88)
sample_1 = rnorm(20,120,15)
Z_stat = (mean(sample_1) - 120) / (15*sqrt(20));Z_stat
```

```
## [1] 0.01776018
```

```
Z_crit = qnorm(0.05);Z_crit
```

```
## [1] -1.644854
```

(b) 100 random samples from the underlying (null) true distribution We find the percentage of 1s to be 0.02, and percentage of 0s to be 0.98. Among the 100 random samples, 2% of them reject the H_0 , 98% of them fail to reject H_0 . This is what we would expect to see since they are sampled from the underlying normal distribution with $\alpha = 0.05$. The Type I error is 2%, which is less than our default $\alpha = 0.05$.

```
set.seed(88)
Z_stat = rep(NA, 100)
decision = rep(NA,100)

for (i in 1:100) {
  Z_stat[i] = (mean(rnorm(20, 120, 15) - 120) / (15/sqrt(20)))
  Z_crit = qnorm(0.05)
  if (Z_stat[i] < Z_crit) {decision[i] = 1}
  else{decision[i] = 0}
}

reject = 0
fail_to_reject = 0

for (i in 1:100) {
  if (decision[i] == 1) (reject = reject + 1)
  if (decision[i] == 0) (fail_to_reject = fail_to_reject + 1)
}

percent_of_reject = reject/100; percent_of_reject
```

```
## [1] 0.06
```

```
percent_of_fail_to_reject = fail_to_reject/100; percent_of_fail_to_reject
```

```
## [1] 0.94
```

(c) 1000 random samples from the underlying (null) true distribution We find the percentage of 1s to be 0.045, and percentage of 0s to be 0.955. Among the 100 random samples, 4.5% of them reject the H_0 , 95.5% of them fail to reject H_0 . This is what we would expect to see since they are sampled from the underlying normal distribution with $\alpha = 0.05$. The Type I error is 4.5%, which is still less than our default $\alpha = 0.05$. However, as we increase the number of random samples, we find the Type I error (0.045) is much more closer to the default $\alpha = 0.05$.

```

set.seed(88)
Z_stat = rep(NA, 1000)
decision = rep(NA, 1000)

for (i in 1:1000) {
  Z_stat[i] = (mean(rnorm(20, 120, 15) - 120) / (15/sqrt(20)))
  Z_crit = qnorm(0.05)
  if (Z_stat[i] < Z_crit) {decision[i] = 1}
  else {decision[i] = 0}
}

reject = 0
fail_to_reject = 0

for (i in 1:1000) {
  if (decision[i] == 1) (reject = reject + 1)
  if (decision[i] == 0) (fail_to_reject = fail_to_reject + 1)
}

percent_of_reject = reject/1000; percent_of_reject

```

```
## [1] 0.045
```

```
percent_of_fail_to_reject = fail_to_reject/1000; percent_of_fail_to_reject
```

```
## [1] 0.955
```

(d) In part(b), we get a Type I error of 0.06, in part (c), we get a Type I error of 0.045. As we increase the number of random samples from the underlying true distribution, we find the Type I error we calculated above are getting closer to our default setting alpha of 0.05.